

## HW5\_MSA 8150

Anutida Sangkla 002602236

3/15/2021

### Question 1

```
Hitters <- read.csv("MyHitters.csv", header=TRUE, sep=",")
str(Hitters)

## 'data.frame':    263 obs. of  20 variables:
## $ AtBat      : int  475 584 484 642 311 281 193 330 625 190 ...
## $ Hits       : int  123 158 127 211 81 76 47 77 179 46 ...
## $ HmRun      : int   27 15 20 14 3 3 10 19 4 2 ...
## $ Runs       : int   76 70 66 107 42 42 21 47 94 24 ...
## $ RBI        : int   93 84 65 59 30 25 29 53 60 8 ...
## $ Walks      : int   72 42 67 52 26 20 24 27 65 15 ...
## $ Years      : int    4 5 7 5 17 8 6 6 5 5 ...
## $ CAtBat     : int  1810 2358 3006 2364 8247 2658 1136 1928 1696 479 ...
## $ CHits      : int   471 636 844 770 2198 657 256 516 476 102 ...
## $ CHmRun     : int   108 58 116 27 100 48 42 90 12 5 ...
## $ CRuns      : int   292 265 436 352 950 324 129 247 216 65 ...
## $ CRBI       : int   343 316 458 230 909 300 139 288 163 23 ...
## $ CWalks     : int   267 134 377 193 690 179 106 161 166 39 ...
## $ League     : int    1 1 1 1 1 0 0 1 0 0 ...
## $ Division   : int    0 0 0 1 1 0 1 1 0 1 ...
## $ PutOuts    : int   226 331 1231 337 153 106 299 149 303 102 ...
## $ Assists    : int   10 20 80 19 223 144 13 8 450 177 ...
## $ Errors     : int    6 4 7 4 10 7 5 6 14 16 ...
## $ Salary     : num  1220 662 1183 740 320 ...
## $ NewLeague  : int    1 1 1 1 1 0 0 1 0 0 ...
```

#### part (a)

```
set.seed(1)
train<- Hitters[1:131,]
test<- Hitters[132:263,]

## fitting a linear regression model
lm.fit<- lm(Salary~., data = train)
summary(lm.fit)

##
## Call:
## lm(formula = Salary ~ ., data = train)
```

```
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -802.73 -178.60   -7.03  123.23 1754.21
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 299.42849   146.63692    2.042  0.04352 *
## AtBat        -2.54027    1.08881   -2.333  0.02145 *
## Hits         8.36682    4.18091    2.001  0.04781 *
## HmRun       11.64512   10.83139    1.075  0.28465
## Runs       -9.09923    5.00294   -1.819  0.07164 .
## RBI         2.44105    4.53052    0.539  0.59110
## Walks       9.23440    3.14657    2.935  0.00406 **
## Years     -22.93673   20.53294   -1.117  0.26638
## CAtBat      -0.18154    0.23637   -0.768  0.44411
## CHits       -0.11598    1.25713   -0.092  0.92666
## CHmRun     -1.33888    2.56414   -0.522  0.60260
## CRuns       3.32838    1.34538    2.474  0.01488 *
## CRBI        0.07536    1.23878    0.061  0.95160
## CWalks     -1.07841    0.67875   -1.589  0.11494
## League     59.76065   134.76740    0.443  0.65831
## Division  -98.86233   66.90363   -1.478  0.14232
## PutOuts     0.34087    0.13298    2.563  0.01171 *
## Assists     0.34165    0.33215    1.029  0.30591
## Errors     -0.64207    6.58517   -0.098  0.92250
## NewLeague  -0.67442   131.05687   -0.005  0.99590
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 350 on 111 degrees of freedom
## Multiple R-squared:  0.5118, Adjusted R-squared:  0.4283
## F-statistic: 6.125 on 19 and 111 DF, p-value: 2.285e-10

## Reporting the test MSE
y_pred <- predict(lm.fit, test)
y_actual<- test$Salary
lm_mse <- (mean((y_actual - y_pred) ^ 2))
sprintf('%s = %10.3f', 'The test MSE', lm_mse)

## [1] "The test MSE = 114780.610"
```

From the result, the test MSE of this model is 114780.610.

#### part (b)

```
set.seed(1)
library(pls)

## Warning: package 'pls' was built under R version 4.0.4
```

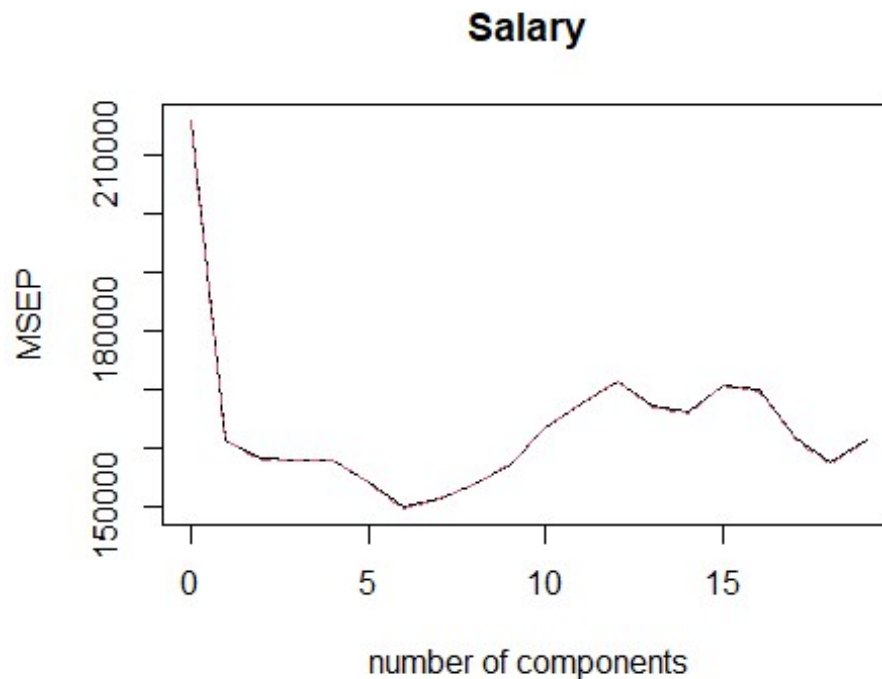
```
##
## Attaching package: 'pls'

## The following object is masked from 'package:stats':
##
##      loadings

pcr.fit <- pcr(Salary~., data = train, scale = TRUE, validation = "LOO")
summary(pcr.fit)

## Data:      X dimension: 131 19
## Y dimension: 131 1
## Fit method: svdpc
## Number of components considered: 19
##
## VALIDATION: RMSEP
## Cross-validated using 131 leave-one-out segments.
##      (Intercept) 1 comps 2 comps 3 comps 4 comps 5 comps 6 comps
## CV      464.6    401.7    397.5    397.5    397.2    392.4    386.9
## adjCV    464.6    401.6    397.4    397.4    397.1    392.4    386.7
##      7 comps 8 comps 9 comps 10 comps 11 comps 12 comps 13 comps
## CV      388.8    392.0    396.4    404.2    409.3    413.9    408.8
## adjCV    388.7    391.9    396.2    404.0    409.0    413.7    408.6
##      14 comps 15 comps 16 comps 17 comps 18 comps 19 comps
## CV      407.4    413.1    412.2    402.0    396.7    401.5
## adjCV    407.1    412.8    411.9    401.7    396.4    401.2
##
## TRAINING: % variance explained
##      1 comps 2 comps 3 comps 4 comps 5 comps 6 comps 7 comps 8
comps
## X      38.89    60.25    70.85    79.06    84.01    88.51    92.61
95.20
## Salary 28.44    31.33    32.53    33.69    36.64    40.28    40.41
41.07
##      9 comps 10 comps 11 comps 12 comps 13 comps 14 comps 15
comps
## X      96.78    97.63    98.27    98.89    99.27    99.56
99.78
## Salary 41.25    41.27    41.41    41.44    43.20    44.24
44.30
##      16 comps 17 comps 18 comps 19 comps
## X      99.91    99.97    100.00    100.00
## Salary 45.50    49.66    51.13    51.18

# The graph of cross-validation in terms of number of components based on
MSEP
validationplot(pcr.fit, val.type = 'MSEP')
```



From the model summary, we will see that the 6 components generate the least cross-validation which equals to 386.9. (We know that this is root mean square error, so if we want MSE, we need to square this quantity.) MSE of 6 components equals to 149691.6.

```
# Test MSE
pcr.pred <- predict(pcr.fit, newdata = test, ncomp = 6)
pcr_mse <- (mean((y_actual - pcr.pred) ^ 2))
sprintf('%s = %10.3f', 'The test MSE', pcr_mse)

## [1] "The test MSE = 96587.921"
```

From the result, the test MSE of this model is 96587.921 which is less than the MSE in part (a). As a result, we can see an improvement in accuracy compared to part (a).

#### part (c)

```
# Center and Scale X
feature <- Hitters[-19]
X <- as.matrix(feature)
X <- scale(X, scale = TRUE)
head(X)

##           AtBat      Hits      HmRun      Runs      RBI      Walks
## [1,]  0.4844122  0.3361993  1.7563137  0.8322203  1.6039009  1.4221315
## [2,]  1.2243624  1.1118170  0.3859982  0.5972930  1.2561785  0.0407924
```

```
## [3,] 0.5455090 0.4248413 0.9569630 0.4406748 0.5220978 1.1919083
## [4,] 1.6180974 2.2863237 0.2718053 2.0460114 0.2902828 0.5012388
## [5,] -0.6289073 -0.5945419 -0.9843172 -0.4990344 -0.8301561 -0.6959218
## [6,] -0.8325634 -0.7053444 -0.9843172 -0.4990344 -1.0233352 -0.9721896
##           Years           CAtBat           CHits           CHmRun           CRuns
CRBI
## [1,] -0.69087451 -0.3706595179 -0.38751381 0.4715523 -0.20900009
0.03890849
## [2,] -0.48226372 -0.1310005959 -0.13296260 -0.1367381 -0.29052218 -
0.04458779
## [3,] -0.06504215 0.1523917061 0.18792619 0.5688787 0.22578439
0.39454083
## [4,] -0.48226372 -0.1283765931 0.07376383 -0.5138782 -0.02783989 -
0.31053893
## [5,] 2.02106574 2.4444581483 2.27678880 0.3742258 1.77772346
1.78923809
## [6,] 0.14356864 0.0001995439 -0.10056518 -0.2583962 -0.11238132 -
0.09406708
##           CWalks           League           Division           PutOuts           Assists           Errors
## [1,] 0.02550157 1.0567429 -1.0172561 -0.2311648 -0.7496555 -0.3925114
## [2,] -0.47817972 1.0567429 -1.0172561 0.1439228 -0.6807283 -0.6952402
## [3,] 0.44208008 1.0567429 -1.0172561 3.3589598 -0.2671650 -0.2411471
## [4,] -0.25474215 1.0567429 0.9792988 0.1653564 -0.6876210 -0.6952402
## [5,] 1.62743530 1.0567429 0.9792988 -0.4919400 0.7184941 0.2129461
## [6,] -0.30776123 -0.9427059 -1.0172561 -0.6598364 0.1739691 -0.2411471
##           NewLeague
## [1,] 1.0730066
## [2,] 1.0730066
## [3,] 1.0730066
## [4,] 1.0730066
## [5,] 1.0730066
## [6,] -0.9284171

# Xtr, Xts, ytr, yts
X.train<- X[1:131,]
X.test<-X[132:263,]
y.tr<- train$Salary
y.ts<- test$Salary
```

Take an SVD of X train to produce the matrices

```
SVD <- svd(X.train)
```

```
# 6 rows and column of ds
d<-diag(SVD$d[1:6])
print(d)
```

```
##           [,1]      [,2]      [,3]      [,4]      [,5]      [,6]
## [1,] 30.52191 0.00000 0.00000 0.00000 0.00000 0.00000
## [2,] 0.00000 23.17414 0.00000 0.00000 0.00000 0.00000
## [3,] 0.00000 0.00000 16.28562 0.00000 0.00000 0.00000
```

```
## [4,] 0.00000 0.00000 0.00000 13.98578 0.00000 0.00000
## [5,] 0.00000 0.00000 0.00000 0.00000 11.07932 0.00000
## [6,] 0.00000 0.00000 0.00000 0.00000 0.00000 10.13017
```

```
dim(d)
```

```
## [1] 6 6
```

```
# 6 columns of U
```

```
u.tr<-SVD$u[,1:6]
```

```
dim(u.tr)
```

```
## [1] 131 6
```

```
# 6 columns of V
```

```
v<-SVD$v[,1:6]
```

```
dim(v)
```

```
## [1] 19 6
```

### Reduction the matrices

```
C<- v %*% solve(d)
```

```
print(C)
```

```
##           [,1]      [,2]      [,3]      [,4]      [,5]
## [1,] -0.0082885758 0.016221829 -0.0014580873 0.002673786 -0.0090218047
## [2,] -0.0081051391 0.016379509 0.0001163945 0.001336259 -0.0063275234
## [3,] -0.0065183483 0.007024864 0.0065529014 -0.022091391 0.0018109076
## [4,] -0.0077135025 0.015959291 0.0039751501 -0.007091922 -0.0033118365
## [5,] -0.0083603520 0.011470960 0.0024224652 -0.010394974 0.0019972247
## [6,] -0.0076375042 0.009278033 -0.0030056134 -0.009460491 -0.0010047132
## [7,] -0.0083825196 -0.012490249 0.0001269619 0.006174271 0.0021290080
## [8,] -0.0103661799 -0.009166026 -0.0026082353 0.007753754 0.0023303944
## [9,] -0.0104200732 -0.008504373 -0.0025134652 0.007030430 0.0029896739
## [10,] -0.0097534225 -0.008481161 0.0031336173 -0.006039137 -0.0003862224
## [11,] -0.0103726261 -0.007761843 -0.0004961119 0.004753712 0.0042329299
## [12,] -0.0107694389 -0.009793481 -0.0010111850 0.001537655 0.0014788721
## [13,] -0.0095559886 -0.008721560 0.0001396867 0.002460232 0.0049494028
## [14,] 0.0008459058 -0.002025148 -0.0399851092 -0.015596816 0.0050121792
## [15,] -0.0001921613 -0.006183547 -0.0006411534 -0.005054658 -0.0885840120
## [16,] -0.0026381464 0.004605511 -0.0078731169 -0.019236195 -0.0029352812
## [17,] -0.0015774937 0.008230119 -0.0143274771 0.044664313 -0.0066815509
## [18,] -0.0015051784 0.007659089 -0.0144548677 0.035825849 -0.0039958207
## [19,] 0.0002052783 -0.001569024 -0.0399372274 -0.015336651 0.0008644338
##           [,6]
## [1,] 0.0018286985
## [2,] 0.0015937031
## [3,] 0.0345520823
## [4,] 0.0067188374
## [5,] 0.0224674920
```

```
## [6,] -0.0241652696
## [7,] 0.0014504660
## [8,] -0.0026643845
## [9,] -0.0044297111
## [10,] 0.0093098095
## [11,] -0.0027227627
## [12,] 0.0001062401
## [13,] -0.0076442291
## [14,] 0.0123780834
## [15,] 0.0040265524
## [16,] -0.0824093964
## [17,] -0.0025945959
## [18,] -0.0019507381
## [19,] 0.0160749330
```

### Fitting a linear model

```
lm_mod<- lm(y.tr~u.tr)
summary(lm_mod)
```

```
##
## Call:
## lm(formula = y.tr ~ u.tr)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -863.64 -172.53  -30.62  120.24 2015.99
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   554.50      32.48   17.073 < 2e-16 ***
## u.tr1        -2866.66     366.71   -7.817 1.99e-12 ***
## u.tr2          648.39     365.80    1.773 0.07877 .
## u.tr3         -619.61     370.29   -1.673 0.09679 .
## u.tr4         -550.47     366.07   -1.504 0.13519
## u.tr5          945.72     365.99    2.584 0.01093 *
## u.tr6        -1020.48     365.91   -2.789 0.00612 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 365.8 on 124 degrees of freedom
## Multiple R-squared:  0.4042, Adjusted R-squared:  0.3753
## F-statistic: 14.02 on 6 and 124 DF,  p-value: 3.941e-12
```

### Testing model

```
Utest<- X.test %*% C
test_mat<- model.matrix(y.ts~Utest)
coefs = coef(lm_mod, id = i)
y_pred = test_mat %*%coefs
```

```

mse <- (mean((y.ts - y_pred) ^ 2))
sprintf('%s = %10.3f', 'The test MSE', mse)

## [1] "The test MSE = 96860.277"

#### Percentage difference
diff<- ((mse - pcr_mse)/pcr_mse)*100
sprintf('%s = %10.3f', 'The difference of MSE between part b and part c',
diff)

## [1] "The difference of MSE between part b and part c = 0.282"

```

From the results, we can see that the MSE of part (c), which is 96860.277, is different than the MSE of part (b), which is 96587.921, up to 1%.

## Question 2

```

auto <- read.csv("Auto.csv", header=TRUE, sep=",")
str(auto)

## 'data.frame': 392 obs. of 8 variables:
## $ mpg : num 14.5 25.5 22.5 13 27.9 18.6 33.5 12 29.5 14 ...
## $ cylinders : int 8 4 6 8 4 6 4 8 4 8 ...
## $ displacement: num 351 140 232 307 156 225 151 383 98 351 ...
## $ horsepower : int 152 89 90 130 105 110 90 180 68 148 ...
## $ weight : int 4215 2755 3085 4098 2800 3620 2556 4955 2135 4657
## $ acceleration: num 12.8 15.8 17.6 14 14.4 18.7 13.2 11.5 16.6 13.5 ...
## $ year : int 76 77 76 72 80 78 79 71 78 75 ...
## $ origin : int 1 1 1 1 1 1 1 1 3 1 ...

```

### part (a)


#### Plotting mpg in terms of horsepower

```

plot(auto$mpg~auto$horsepower, xlab = 'Horsepower', ylab = 'mpg')

```





### Fitting a polynomial model

```
poly.fit <- lm(mpg~poly(horsepower,6), data = auto)
summary(poly.fit)
```

```
##
## Call:
## lm(formula = mpg ~ poly(horsepower, 6), data = auto)
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
	-15.595	-2.571	-0.269	2.209	15.362

```
##
## Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	23.4459	0.2177	107.715	< 2e-16 ***
poly(horsepower, 6)1	-120.1377	4.3096	-27.877	< 2e-16 ***
poly(horsepower, 6)2	44.0895	4.3096	10.231	< 2e-16 ***
poly(horsepower, 6)3	-3.9488	4.3096	-0.916	0.36008
poly(horsepower, 6)4	-5.1878	4.3096	-1.204	0.22941
poly(horsepower, 6)5	13.2722	4.3096	3.080	0.00222 **
poly(horsepower, 6)6	-8.5462	4.3096	-1.983	0.04807 *


```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.31 on 385 degrees of freedom
```

```
## Multiple R-squared:  0.6998, Adjusted R-squared:  0.6951  
## F-statistic: 149.6 on 6 and 385 DF,  p-value: < 2.2e-16
```

### Plot 95% confidence interval

```
## Plotting 95% confidence interval  
hpwlims<- range(auto$horsepower)  
hpw.grid<- seq(from = hpwlims[1], to = hpwlims[2])  
pred<- predict(poly.fit, newdata = list(horsepower = hpw.grid), interval =  
'confidence', level = 0.95)  
plot(auto$horsepower, auto$mpg, xlim = hpwlims, cex = 0.5, col = 'darkgrey',  
xlab = 'Horsepower', ylab = 'mpg', main = "95% Confidence Interval of Degree-  
6 Polynomial")  
lines(hpw.grid, pred[,1], lwd = 2, col = 'lightblue')  
matlines(hpw.grid, pred[,-1], lwd = 1, col= 'darkblue', lty = 2)
```

### 95% Confidence Interval of Degree-6 Polynomial



## part (b)

### Fitting model to a natural spline

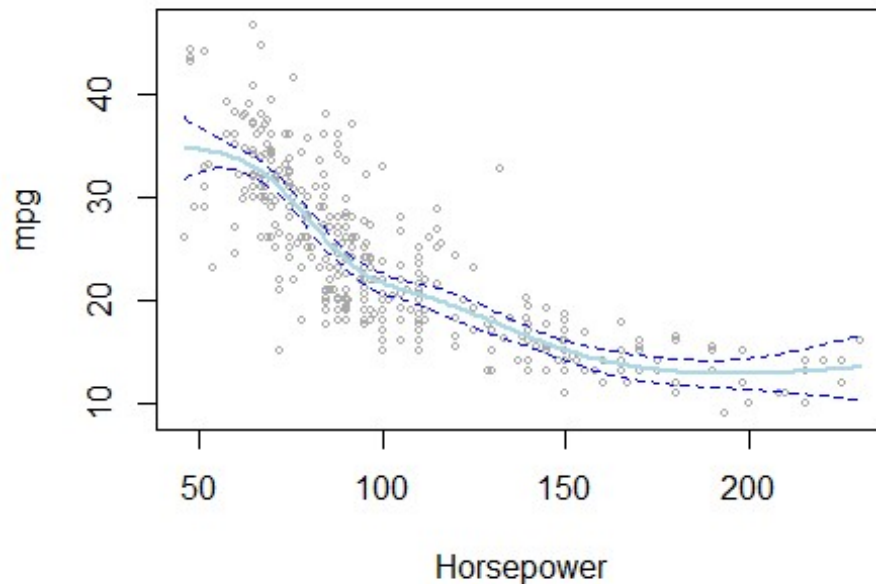
```
library(splines)
fit<- lm(mpg~ns(horsepower, 6), data = auto)
summary(fit)

##
## Call:
## lm(formula = mpg ~ ns(horsepower, 6), data = auto)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -15.9491  -2.6183  -0.1595   2.3508  15.1349
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      34.738      1.509   23.021 < 2e-16 ***
## ns(horsepower, 6)1  -8.210      1.594   -5.149 4.18e-07 ***
## ns(horsepower, 6)2 -13.046      1.835   -7.108 5.76e-12 ***
## ns(horsepower, 6)3 -14.577      1.886   -7.730 9.50e-14 ***
## ns(horsepower, 6)4 -22.802      1.624  -14.039 < 2e-16 ***
## ns(horsepower, 6)5 -22.758      3.512   -6.480 2.81e-10 ***
## ns(horsepower, 6)6 -20.849      1.742  -11.967 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.302 on 385 degrees of freedom
## Multiple R-squared:  0.7009, Adjusted R-squared:  0.6962
## F-statistic: 150.4 on 6 and 385 DF,  p-value: < 2.2e-16
```

### Plotting 95% confidence interval

```
hpwlms<- range(auto$horsepower)
hpw.grid<- seq(from = hpwlms[1], to = hpwlms[2])
pred<- predict(fit, newdata = list(horsepower = hpw.grid), interval =
'confidence', level =0.95)
plot(auto$horsepower, auto$mpg, xlim = hpwlms, cex = 0.5, col = 'darkgrey',
xlab = 'Horsepower', ylab = 'mpg', main = "95% Confidence Interval of Degree-
6 Spline ")
lines(hpw.grid, pred[,1], lwd = 2, col = 'lightblue')
matlines(hpw.grid, pred[,-1], lwd = 1, col= 'blue', lty = 2)
```

## 95% Confidence Interval of Degree-6 Spline



From the result, natural spline fit in part (b) seems to have a narrower confidence interval around the boundaries compared to part (a) when the horsepower is large. However, part (a) seems to have a narrower confidence interval around the boundaries compared to part (b) when the horsepower is small.

### part (c)

```
train<- auto[1:350,]  
test<- auto[351:392,]
```

#### ## Fitting linear model

```
lm.fit<- lm(mpg~horsepower+acceleration+year, data = train)  
summary(lm.fit)
```

```
##
```

```
## Call:
```

```
## lm(formula = mpg ~ horsepower + acceleration + year, data = train)
```

```
##
```

```
## Residuals:
```

```
##      Min       1Q   Median       3Q      Max  
## -11.7704  -3.0253  -0.6865   2.1362  15.4392
```

```
##
```

```
## Coefficients:
```

```
##              Estimate Std. Error t value Pr(>|t|)  
## (Intercept)  -0.641629   6.015622  -0.107    0.915  
## horsepower   -0.160214   0.008482 -18.888 < 2e-16 ***
```

```

## acceleration -0.581611    0.112584   -5.166 4.05e-07 ***
## year          0.657839    0.068474    9.607 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.301 on 346 degrees of freedom
## Multiple R-squared:  0.7079, Adjusted R-squared:  0.7054
## F-statistic: 279.5 on 3 and 346 DF,  p-value: < 2.2e-16

## Testing the model
y_pred <- predict(lm.fit, test)
y_actual<- test$mpg
lm_mse <- (mean((y_actual - y_pred) ^ 2))
sprintf('%s = %10.3f', 'The test MSE', lm_mse)

## [1] "The test MSE =      12.059"

## Fitting GAM model
library(splines)
library(gam)

## Warning: package 'gam' was built under R version 4.0.4

## Loading required package: foreach

## Warning: package 'foreach' was built under R version 4.0.3

## Loaded gam 1.20

gam.fit<-gam(mpg~ns(horsepower, 4)+ ns(acceleration,4)+ year, data = train)
summary(gam.fit)

##
## Call: gam(formula = mpg ~ ns(horsepower, 4) + ns(acceleration, 4) +
##      year, data = train)
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -11.718  -1.935  -0.272   1.480  13.055
##
## (Dispersion Parameter for gaussian family taken to be 11.6243)
##
##      Null Deviance: 21912.78 on 349 degrees of freedom
## Residual Deviance: 3952.262 on 340 degrees of freedom
## AIC: 1863.696
##
## Number of Local Scoring Iterations: 2
##
## Anova for Parametric Effects
##
##              Df  Sum Sq Mean Sq F value    Pr(>F)
## ns(horsepower, 4)    4 15282.5   3820.6  328.676 < 2.2e-16 ***

```

```
## ns(acceleration, 4)    4    740.1    185.0    15.917 5.913e-12 ***
## year                   1    1937.9    1937.9    166.714 < 2.2e-16 ***
## Residuals              340    3952.3     11.6
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

## Testing the model
y_pred <- predict(gam.fit, test)
y_actual<- test$mpg
gam_mse <- (mean((y_actual - y_pred) ^ 2))
sprintf('%s = %10.3f', 'The test MSE', gam_mse)

## [1] "The test MSE =      8.100"
```

### Comparing two models:

```
library(dplyr)

##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union

data.frame(Models = c("Linear Regression", "GAM"),
           Test_Error = c(lm_mse, gam_mse))

##           Models Test_Error
## 1 Linear Regression 12.059045
## 2              GAM    8.100448
```

From the result above, the test error of GAM model, which equals 8.10, is smaller than the test error of linear regression model.